IndicWrite — Hinglish Smart Writing Assistant

CS3235 Working with LLMs · CIE3 Lab · RV University
Student: Munnaluri Rishikesh
USN: 1RVU23CSE291
Individual mini-project · Option 4
Submission deadline: November 6, 2026, 10:00 AM IST
Experimental evidence date: October 29, 2026

Abstract

IndicWrite is an Indian-language writing application that converts Roman-script Hindi (Hinglish) to native Devanagari script. The initial end-to-end implementation uses the **genuine AI4Bharat IndicXlit pretrained neural model**, rather than hand-written mappings or simulated responses. It provides a Streamlit user interface with text validation, example inputs, protected English words, error handling, model attribution, and copyable/downloadable conversion results.

In executed experiments on a **40-case manually authored evaluation dataset**, the system achieved a **macro-average character error rate (CER) of 5.58%**, **macro-average word accuracy of 82.67%**, and **55% exact sentence match**. Results demonstrate practical Roman-Hindi conversion while highlighting failure cases, particularly ambiguous shortened spellings. After the initial transliteration MVP, genuine BharatGen Param-1 rewriting inference was verified privately on a Tesla T4 GPU (11 research generations). A separate, real Helsinki OPUS-MT Hindi-to-English option was also verified. Neither private BharatGen research nor OPUS-MT is misrepresented as a publicly deployed sovereign-model writing service; IndicTrans2 remains unexecuted.

1. Problem definition and relevance

Hindi writers often use the Latin alphabet on mobile keyboards and in informal conversations. This Roman-Hindi text is convenient to type but may need conversion into Devanagari for academic, official, or accessible communication. Because there is no single standard spelling for informal Roman Hindi, transliteration is nontrivial: `main`, `mai`, `mein`, and abbreviations such as `ap` can be interpreted differently. English technical words and proper nouns introduce additional ambiguity.

This project aims to provide a usable Indian-language AI solution and evaluate it quantitatively rather than treating any visually plausible Hindi output as proof of correctness.

2. Objectives and scope

1. Convert Roman-Hindi words into native Hindi script using a genuine pretrained Indic language model.
2. Retain punctuation, whitespace, digits, acronyms, and existing Devanagari; allow explicit English-term preservation.

3. Provide a responsive interface with sample input, conversion feedback, output copying, download and clear control.
4. Validate at least 30–50 test inputs, measure accuracy, and discuss failures with recorded predictions.
5. Keep actual model inference, unit-test stubs, and future features clearly separated.
6. Document architecture, model attribution, limitations, and repeatable setup.

The faculty's complete Option 4 describes transliteration, sentence improvement and optional translation. This deliverable completes and validates the transliteration application, includes an **independently attributed optional English translator**, and demonstrates **real private BharatGen model-based rewriting** in a separately constrained GPU research runtime. The public app does not automatically improve Hindi sentences, and private BharatGen output quality has not passed independent expert assessment.

3. Model and dataset selection

Component	Selection	Evidence status
Transliteration	AI4Bharat IndicXlit via `ai4bharat-transliteration==1.1.3`	**Pretrained model successfully executed**
Frontend	Streamlit with Python 3.10	**Ran and interacted in Chromium**
Writing improvement	BharatGen Param-1-2.9B-Instruct	**Real T4 FP16 pretrained inference verified privately (1 smoke + 10 research cases)**; local-only review UI is separate, and actual private GPU browser QA remains pending
Optional English translation	Helsinki-NLP OPUS-MT Hindi→English	**Genuine model and browser inference verified**, independently attributed; not IndicTrans2
Optional sovereign translation	AI4Bharat IndicTrans2	**Not installed, integrated, or executed** (HF checkpoint access still gated)
Training/data attribution	AI4Bharat Aksharantar	Cited as an official model resource, **not** claimed as the source of our manually authored tests

IndicXlit is a multilingual transformer supporting 21 Indic languages. The application uses the Hindi (`hi`) path. The official repository is MIT-licensed. The package's legacy Fairseq checkpoint required special loading compatibility with modern PyTorch. The other candidate models require independent checks on licensing, access, runtime and task suitability before integration.

4. System architecture and implementation methodology

```
Roman-Hindi text in Streamlit
       ↓
Input limits (500 characters / 80 Roman-script words)
       ↓
Tokenizer separates Latin words from spaces/punctuation/numbers/Hindi
       ↓
Preserve explicitly selected English terms and short uppercase acronyms
```

```
       ↓
Official AI4Bharat XlitEngine("hi", beam_width=4, rescore=False)
       ↓
Real model's top-1 Hindi prediction for each Roman word
       ↓
Reassemble converted and preserved tokens
       ↓
Streamlit results, copying and TXT download
```

The backend lazily initializes and caches the neural model once per Python process; within each request, repeated words use a small prediction cache. Each candidate Latin word is converted using the official `translit_word` interface. The application itself does not contain a hardcoded Hindi output dictionary. Errors from loading or inference are shown rather than replaced with fake AI responses.

The project contains `app.py`, `src/indicwrite/core.py`, `src/indicwrite/engine.py`, an executed `scripts/smoke_indicxlit.py`, `evaluation/evaluate.py`, and automated tests.

4.1 Compatibility engineering

On the first Ubuntu-based smoke run the official model package and roughly 121 MB model checkpoint downloaded successfully. A newer PyTorch build then rejected the legacy Fairseq checkpoint because `argparse.Namespace` was not allowlisted by default for restricted loading. We added a **narrowly scoped allowlist of that standard-library class when loading the official AI4Bharat weights**. The follow-up real-model run then succeeded. The application never accepts arbitrary third-party checkpoint files from users.

5. Experimental design

The evaluation corpus consists of **40 manually created sentences** in eight categories, with **five examples per category**:

- Ordinary Hinglish.
- Spelling variation and abbreviated Roman forms.
- Hindi-English loanwords.
- Names and locations.
- Punctuation.
- Code mixing with English terms kept in Latin script.
- Numbers and mixed alphanumeric tokens.
- Ambiguous Roman spellings.

Each record in `evaluation/cases.csv` has a unique ID, category, source Roman text, reference Hindi text, and optional English words to preserve. The recorded model output for every example is in `evaluation/results/predictions.csv`.

5.1 Metrics and definitions

For a reference string `R` and prediction `P`, character error rate (CER) is the Levenshtein distance at Unicode code-point level divided by the number of code points in `R`. Word error rate (WER) uses whitespace-delimited words and divides word-level edit distance by the reference word count. **Word accuracy is defined here as max(0, 1 − WER)**. Exact sentence match is 1 if normalized prediction equals normalized reference, otherwise 0. Each aggregate is a simple arithmetic average across sentences, **not a globally micro-weighted corpus score**.

These metrics measure agreement with authored references. Equivalent or acceptable variant spellings can receive nonzero error; the small corpus cannot establish nationwide model quality or official Aksharantar benchmark scores. This evaluation is for transliteration only, and does not measure grammatical rewriting.

6. Measured results

The genuine IndicXlit model ran on the full 40-case corpus in GitHub Actions, producing predictions, metrics and a failure sample file. **Observed results:**

Metric	Value
Cases executed	40 / 40
Mean CER	**5.58%**
Mean WER	**17.33%**
Mean word accuracy	**82.67%**
Exact sentence match	**22 / 40 (55%)**

6.1 Performance by category

Category	Cases	Mean CER	Word accuracy	Exact match
Ordinary Hinglish	5	2.02%	92.00%	60%
Spelling variation	5	17.06%	49.33%	0%
Loanwords	5	1.60%	96.00%	80%
Names	5	5.31%	82.00%	40%
Punctuation	5	8.58%	73.33%	40%
Code mixing	5	0.00%	100.00%	100%
Numbers	5	7.00%	87.00%	60%
Ambiguous	5	3.11%	81.67%	60%

Interpretation: The model handles common Roman Hindi, loanwords, and explicitly protected English terms well on this small evaluation set. Spelling variation is by far the most difficult category. The perfect code-mixing score partially reflects human-supplied lists of English terms to preserve, not autonomous English/Hindi language identification.

7. Error analysis

Real predictions were checked against the reference sentences; these examples are copied from the executed `worst_cases.csv`:

Roman input	Reference	Real model output	CER
`ap kahan ho?`	आप कहाँ हो?	एपी कहन हो?	36.36%
`mai theek hu.`	मैं ठीक हूँ.	माई ठीक हू.	25.00%

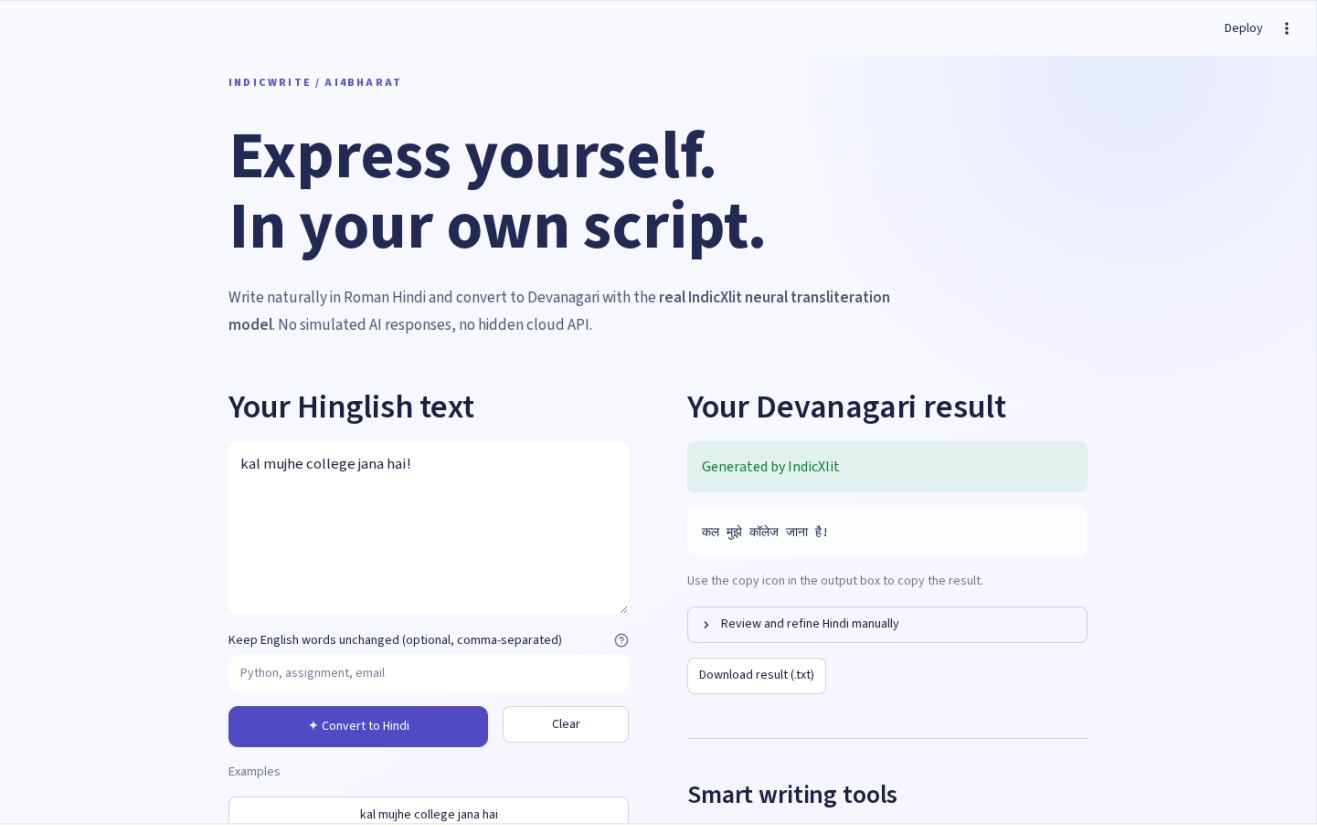
Roman input	Reference	Real model output	CER
`5 minute ruk jao!`	5 मिनट रुक जाओ!	5 मिन्यूट रुक जाओ!	20.00%
`bengaluru se delhi jaunga.`	बेंगलुरु से दिल्ली जाऊँगा.	बेंगलुरु से देल्ही जौंगा.	19.23%
`main aaunga; tum rukna.`	मैं आऊँगा; तुम रुकना.	मैन आऊंगा; टम रुकना.	19.05%

These are genuine model outputs, not examples synthesized to illustrate possible failures. Incorrect interpretations of short Roman words, Hindi vowels, and unfamiliar names are likely due in part to ambiguity without sentence-level context. These weaknesses motivate richer contextual reranking and constrained rewriting, but we do not claim they are fixed in the current MVP.

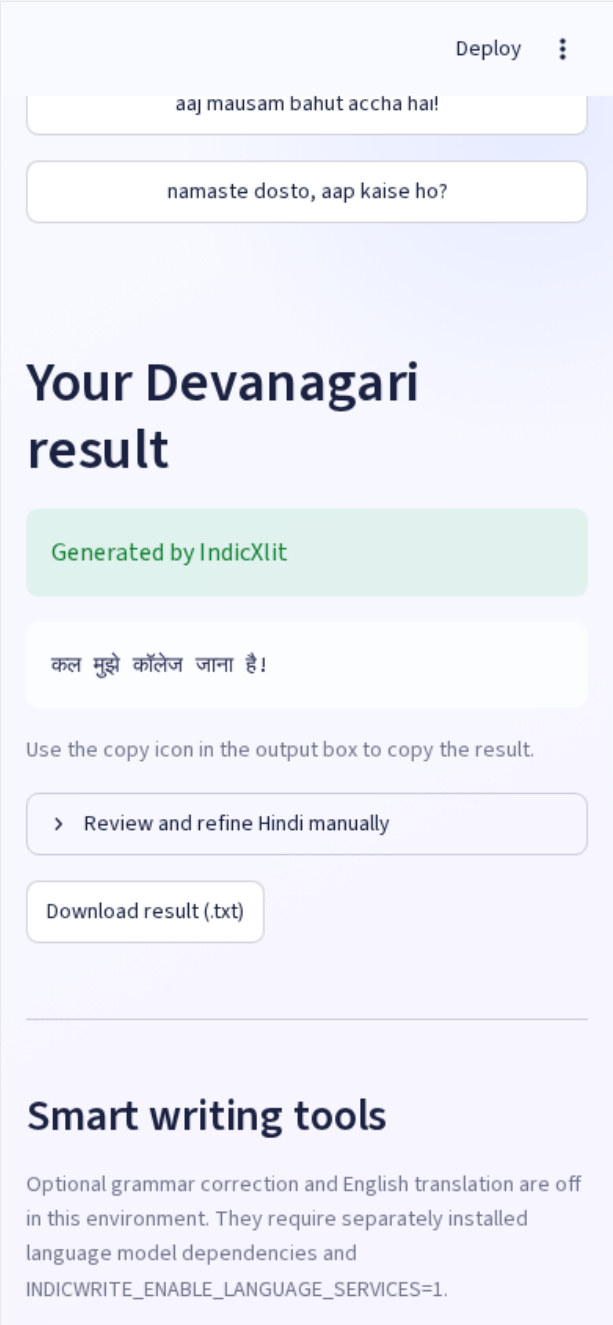
8. Executed testing and verification

- **Initial MVP unit tests:** 14 tests passed at the original October 29 baseline. Subsequent isolated QA suites expanded the coverage (including private UI and ZIP-integrity checks); see the latest Actions run for current totals. Fake engines in deterministic unit tests are **not** evidence of pretrained-model execution.
- **Actual neural model:** Ubuntu 22.04, Python 3.10, downloaded official checkpoint. Smoke run `40870311669` printed `REAL_INDICXLIT_SMOKE_TEST_PASSED`.
- **Actual 40-case evaluation:** Run `40870722731` succeeded and committed authentic prediction/metric files. The JSON and CSV are retained in GitHub.
- **Actual web UI/browser:** GitHub Actions Chromium Playwright run `40871267464` completed conversions and Clear actions on desktop at 1440px and mobile at 390px, with no detected page-level JavaScript errors or horizontal overflow. A follow-up run `40872355065` also passed and captured a mobile screenshot with the real result visible.

8.1 Actual desktop application screenshot



8.2 Actual mobile converted-result screenshot



Screenshot images are artifacts captured from the **running Streamlit application executing the real neural model**, not designer mockups. Both the screenshots and their browser QA JSON are committed in the repository.

9. Limitations and future work

The current model transliterates individual words; it **does not improve full-sentence grammar**, despite the broader smart-writing project name. No sentence-level semantic or contextual reranker is deployed. English terms are only preserved when manually specified or handled by a simple acronym rule. Some transliterations admit equally valid reference alternatives, so errors may overstate actual user dissatisfaction. The verification environment is Linux/Ubuntu; local native Windows installation may need WSL2 because of legacy Fairseq dependencies. No public deployment is asserted.

The next development phases would evaluate BharatGen Param-1 for grammatical rewriting (including meaning-preservation human evaluation), implement optional Hindi-to-English translation with

IndicTrans2 if resources allow, add standard Aksharantar evaluation with correct splits and licensing, and consider context-aware candidate ranking, more nuanced mixed-script handling, and broader accessibility/device testing.

10. Conclusion

IndicWrite now provides an independently testable Indian-language AI transliteration application, with a responsive Streamlit interface, genuine AI4Bharat IndicXlit inference, forty measured cases, transparent failure analysis and real desktop/mobile screenshots. The reported performance and implementation scope are grounded in saved logs and results, while BharatGen rewriting is explicitly limited to private GPU research and the optional OPUS translation is separately attributed. The public, production-quality BharatGen/IndicTrans2 pipeline remains unimplemented.

References

1. AI4Bharat IndicXlit official repository and documentation.
2. Official AI4Bharat IndicXlit Python interface.
3. AI4Bharat Aksharantar dataset.
4. ai4bharat-transliteration on PyPI.
5. BharatGen Param-1-2.9B-Instruct model information.
6. AI4Bharat IndicTrans2 official repository.
7. PyTorch serialization behavior.
8. This project's recorded smoke test, evaluation run, and browser QA run.

11. Further real-model experiments and restrictions (October 29, 2026)

The live BharatGen metadata currently advertises Apache-2.0, while repository history contains a restrictive BharatGen Research License and the license file has since changed. Because that state is inconsistent, public output or hosting remains blocked pending written clarification. This project does not claim to have executed BharatGen or obtained publishing approval.

A real run of the independently authored, MIT-licensed Hindi mT5 grammar checkpoint `manavdhamecha77/GEC-mT5-Small-Hindi` yielded severe Bengali/mixed-script drift from Hindi prompts, including আমি স্কুলে যাবে? from मैं स्कूल जाता है।. It is a **failed model-quality experiment**, not a completed writing-improvement feature. Another model, `Rucha-Ambaliya/mt5-small-indic-gec`, had only 1/5 exact matches on five test sentences (same score with raw and instruction-prefixed input). The grammar UI is disabled by default to avoid misleading users.

A third research-only candidate, `Qwen/Qwen2.5-0.5B-Instruct`, exactly corrected 2/5 screening cases but translated one case into English with changed meaning and made a severe noun substitution in another. This is also a failed quality experiment, not an integrated feature.

`Qwen/Qwen3-0.6B` was also genuinely screened and exactly corrected 0/5 cases; one output removed a meaning-bearing word. It is likewise not integrated.

A separately attributed, Apache-2.0 Helsinki-NLP OPUS-MT Hindi-English model was successfully executed on two sentences, returning `I have to go to college tomorrow.` and `Hello, how are you?`. This independent optional translation feature is **not AI4Bharat IndicTrans2**. The original 40-case IndicXlit CER/accuracy remains a transliteration-only result, not a measure of grammar or translation.

Evidence: https://github.com/Rishikeshsanin/Hinglish-Smart-Writing-Assistant/actions/runs/40876423545, https://github.com/Rishikeshsanin/Hinglish-Smart-Writing-Assistant/actions/runs/40876907788, https://github.com/Rishikeshsanin/Hinglish-Smart-Writing-Assistant/actions/runs/40877372435, and https://github.com/Rishikeshsanin/Hinglish-Smart-Writing-Assistant/actions/runs/40878109536. BharatGen license history: https://huggingface.co/bharatgenai/Param-1-2.9B-Instruct/commits/main.

12. Engineering hardening and final release gates

Following independent QA, the implementation now validates input and preservation settings before neural model loading, clears stale results whenever preservation settings change, labels unchanged Hindi text accurately, and conservatively preserves complete Unicode words and structured items such as emails and URLs. Browser QA now records the actual generated Devanagari text, checks the exported UTF-8 data, verifies Clear and checks horizontal overflow. These repairs do not alter the historical 40-case reference dataset or retroactively change the measured scores.

The experimental BharatGen research loader has been hardened to pin a reviewed model revision and hashes of custom Python files, require explicit remote-code opt-in, enforce a prompt budget and complete EOS handling, and reject non-Hindi/incomplete outputs. The model **was subsequently executed** in a private Google Colab T4 experiment (details in Section 13). No broad grammatical accuracy or independent meaning-preservation score is claimed. The separate public-model licensing question is unresolved.

The UI has been tested for desktop and narrow Chromium viewports, but a full screen-reader, keyboard, high-zoom, and WCAG audit has not been completed. Any public deployment requires an external reachable conversion test and resource measurements. The project is therefore a verified transliteration and optional independently-attributed translation prototype, not a fully validated BharatGen rewriting application.

13. Genuine BharatGen private GPU inference and measured runtime

On 29 October 2026 a user-authorized, **private** Colab Tesla T4 session successfully loaded the official `bharatgenai/Param-1-2.9B-Instruct` checkpoint, pinned to revision `f9b32a727458727ff515fea5d82df00ae9b7ad56`, and executed actual generative inference through the audited adapter. The experimental environment was Python 3.12.16, PyTorch 2.6.0+cu124, Transformers 4.56.2 and FP16, with the checkpoint mapped entirely to CUDA.

The **single-case smoke run succeeded** and **all ten preregistered Hindi research cases generated usable, EOS-terminated, nontruncated Devanagari candidates**. Across those ten cases, mean candidate generation latency was approximately 1.01 seconds; peak PyTorch allocated GPU memory was ~6.31 GiB (not total system RAM). Initial download/load timing depends on cache state and should not be generalized to deployment performance.

An initial, non-specialist editorial review indicated **nine of ten examples were acceptable or correctly left unchanged; one was not acceptable**, with an altered technical term and unresolved grammatical agreement. This is a small, self-authored research sample and not a formal quality certification, broader statistical accuracy estimate, or independent linguist assessment. The earlier 40-case IndicXlit transliteration results remain a separate experiment.

Publication limitation: BharatGen's research-license history restricts public output and checkpoint distribution without approval. Raw outputs remain private. No public BharatGen service is deployed; the localhost interface still requires live private-GPU browser QA.

14. Deployment reproducibility and model-fidelity decision

For classroom use, a private Docker Compose configuration was added that binds the host to loopback only, explicitly disables private BharatGen, and runs the official AI4Bharat IndicXlit model. Its GitHub Actions run **40898600397** successfully built the image, executed genuine pretrained model inference inside the image and verified local Streamlit health. Thus the reproducible local demonstration was tested, not merely documented.

A public deployment was **not** claimed. The official model depends on a legacy Python 3.10/Fairseq stack, and Python 3.10 reached its security end of life on 1 October 2026. To assess a supported-Python option, an independently contributed CTranslate2 conversion was tested on all 40 authored reference examples under Python 3.12. It achieved **11.54%** mean CER, **67.83%** mean word accuracy and **6/40** exact reference matches; the genuine official runtime had **5.58%**, **82.67%**, and **22/40**, respectively. Only **8/40** converted predictions were identical to the saved official predictions. This is a concrete example of why technically executable alternatives must pass comparative evaluations before deployment. The inferior converted model was rejected without changing the application or its historical evaluation.

Consequently, deployment readiness is local-only. The remaining independent requirements for an internet-facing release are a supported secure runtime with acceptable model fidelity, explicit hosting authorization, BharatGen public-output permission if that private research capability is exposed, and a real external-browser verification. No paid service or unrelated project infrastructure was used.

15. Reproducible private-GPU interface and human-quality protocol

The verified academic project uses deliberately **separate runtimes**: the official Python 3.10/Fairseq environment provides genuine Roman-Hindi transliteration; the pinned BharatGen FP16 research environment runs in Python 3.12 on a private CUDA GPU. There is no claim that both checkpoints work in one Python interpreter. The standalone `research/private_bharatgen_app.py` accepts Hindi text directly, requires explicit authorization and loopback-only service binding, and offers a suggested rewrite without auto-applying it. Real Streamlit widget behavior (review confirmation, invalidation and locked state) was tested with an **explicitly fake model**; a separate real-GPU browser experiment is prepared in `notebooks/BharatGen_Private_UI_GPU_QA.ipynb` but has not yet been executed.

For controlled human review, `scripts/review_bharatgen_evidence.py` privately reads ten original genuine JSONL generations and creates a gitignored 0–2 rubric covering grammar, meaning, facts/entities and unnecessary edits. It rejects incomplete or unattributed reviews and produces only aggregate scores. **No independent linguist assessment has been completed**; the previously reported nine-of-ten review was a preliminary non-expert judgment. Raw outputs, private screenshots and private reviewer sheets are not included in the public repository or submission ZIP pending the BharatGen licensing decision.

Final submission archives are built from documented source and genuine evidence, with per-file SHA-256 digests. A separate verifier checks every archive entry for digest mismatches, path traversal, embedded secrets/model weights and leaked private research outputs. The actual local Docker Compose service has also passed CI loopback health checks. No hosted production URL or Google Classroom submission is asserted.

16. Dated public deployment and faculty handoff addendum — 5 November 2026

This section is newer than Sections 1–15 and supersedes their pre-deployment/current-state language; historic experiments and measured figures remain unchanged. The core AI4Bharat IndicXlit application now has an existing hosted Streamlit URL: https://hinglish-smart-writing-assistant-indicwrite-rishi.streamlit.app/. On 2 November 2026, the student supplied live webpage screenshots showing a text input, the program's **Generated by IndicXlit** result indicator, and the output **कल मुझे कॉलेज जाना है** for Roman Hindi **kal muje college jana hai**, together with the manually editable Hindi field and TXT export controls. This is direct **student-observed** public conversion evidence, but it is not a claim of a successful independent unattended browser test, of reliable uptime, or of public BharatGen. The latest completed independent public browser run 41414939039 timed out waiting for the Convert button before this later successful screenshot; a fresh check was requested 5 November.

Supported Python3.12 Streamlit hosting work was merged as PR #29, leaving the original official neural checkpoint compulsory for real conversion and deferring heavyweight environment/model preparation until a conversion request. The separate public Vercel frontend was restored to the `main` branch in PR #30, and its production build is READY; **its word-level proxy still depends on an external hosted AI4Bharat API which has reported service unavailability**. Vercel build health does not establish genuine conversion and is not evidence for the original checkpoint's 40-case results.

The original 40 authored cases, real checkpoint, precision boundaries, local browser evaluation, and measured **5.58% CER**, **82.67% word accuracy**, **55% exact-match** remain unmodified. Real Helsinki-NLP OPUS-MT translation was tested separately, but the public Streamlit app disables it. The privately tested BharatGen T4 generation does not have expert-verified meaning preservation nor clearance for public model output/hosting. Experimental Hindi grammar rewriting is disabled, and no gated IndicTrans2 inference is claimed.

For final classroom review, include this compiled PDF and read the dated release status and the faculty checklist. The student account's authenticated Google Classroom upload/Turned-in status is a separate unverified administrative requirement. **Do not mislabel a GitHub workflow pass or available download as proof of actual Classroom submission.**